

MSU CSE 803 Computer Vision: Homework 5

Instructions

- This homework is **due at 11:59:59 p.m. on Wednesday December 1st, 2021.**
- The submission includes two parts:
 1. A pdf file as your write-up, including your plots and answers to all the questions and key choices you made.
The write-up must be an electronic version. **No handwriting, including plotting questions.** \LaTeX is recommended but not mandatory.
You might like to combine several files to make a submission. Here is an example online link for combining multiple PDF files: <https://combinepdf.com/>.
 2. A zip file including all your code, and files specified in questions with Submit, all under the same directory. You can submit Python code in either .py or .ipynb format.

Python Environment

We are using Python 3.7 for this course. You can find references for the Python standard library here: <https://docs.python.org/3.7/library/index.html>. To make your life easier, we **recommend** you to install Anaconda 5.2 for Python 3.7.x (<https://www.anaconda.com/download/>). This is a Python package manager that includes most of the modules you need for this course.

We will make use of the following packages extensively in this course:

- Numpy (<https://docs.scipy.org/doc/numpy-dev/user/quickstart.html>)
- OpenCV (<https://opencv.org/>)
- PyTorch (<https://pytorch.org/>)
- Matplotlib (http://matplotlib.org/users/pyplot_tutorial.html)

In this homework, you are required to use PyTorch for building and training neural networks. Install PyTorch as `torch` and `torchvision` for datasets. You will also need `matplotlib.pyplot` to visualize results and `tqdm` to display a progress bar.

As a deep learning library, PyTorch performs backpropagation automatically for you and trains your network faster. To familiarize yourself with PyTorch, you can refer to the Discussion 10 notebook at <https://piazza.com/class/jyxc5yj1x6t3wa?cid=35> which contains a brief introduction. For this homework we have also provided some starter code with comments.

1 Fashion-MNIST Classification [35 pts]

In this part, you will implement and train Convolutional Neural Networks (ConvNets) in **PyTorch** to classify images. Unlike HW4, backpropagation is automatically inferred by PyTorch, so you only need to write code for the forward pass.

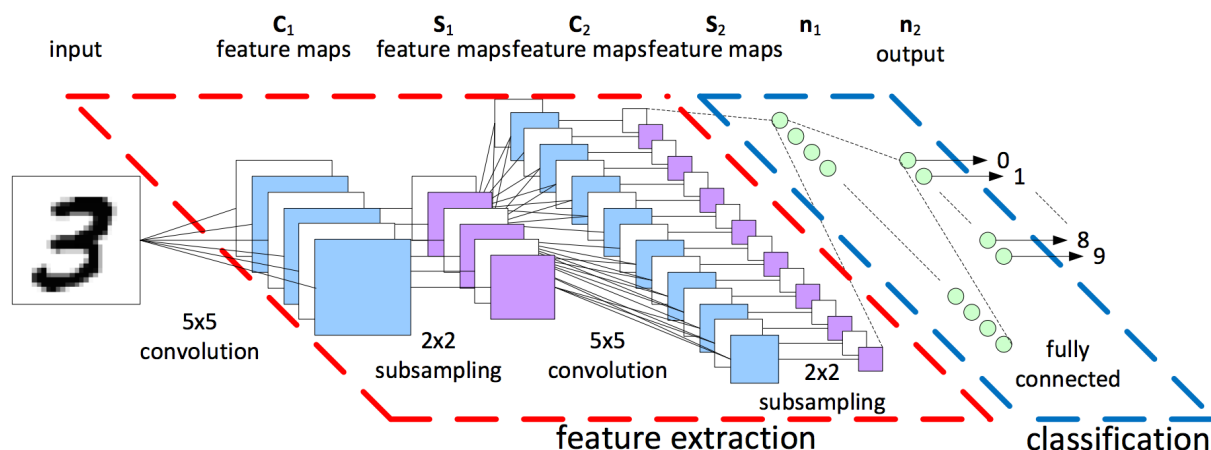


Figure 1: Convolutional Neural Networks for Image Classification¹



Figure 2: Example images from Fashion MNIST dataset [2]

The dataset we use is Fashion-MNIST dataset, which is available at <https://github.com/zalandoresearch/fashion-mnist> and in `torchvision.datasets`. Fashion-MNIST has 10 classes, 60000 training+validation images (we have splitted it to have 50000 training images and 10000 validation images, but you can change the numbers), and 10000 test images. We have provided some starter code in `part1.py` where you need to modify and experiment with the following:

¹Image comes from <https://medium.com/data-science-group-iitr/building-a-convolutional-neural-network-in-python-with-tensorflow-d251c3ca8117>

- The architecture of the network (define layers and implement forward pass)
- The optimizer (SGD, RMSProp, Adam, etc.) and its parameters. (`weight_decay` is the L2 regularization strength)
- Training parameters (batch size and number of epochs)

You should train your network on training set and change those listed above based on evaluation on the validation set. You should run evaluation on the test set **only once** at the end.

Complete the following:

1. **Submit** a program which trains with your best combination of model architecture, optimizer and training parameters, and evaluates on the test set to report an accuracy at the end. (15 pts)
2. **Report** the detailed architecture of your best model. Include information on hyperparameters chosen for training and a plot showing both training and validation loss across iterations. (10 pts)
3. **Report** the accuracy of your best model on the test set. We expect you to achieve over **90%**. (10 pts)

Hints: Read PyTorch documentation for `torch.nn` at <https://pytorch.org/docs/stable/nn.html> and pick layers for your network. Some common choices are

- `nn.Linear`
- `nn.Conv2d`, try different number of filters (`out_channels`) and size of filters (`kernel_size`)
- `nn.ReLU`, which provides non-linearity between layers
- `nn.MaxPool2d` and `nn.AvgPool2d`, two kinds of pooling layer
- `nn.Dropout`, which helps reduce overfitting

Your network does not need to be complicated. We achieved over 90% test accuracy with two convolutional layers, and it took less than 15 mins to train on a laptop without GPU. You will get partial credits for any accuracy over **50%**, so do not worry too much and spend your time wisely.

2 Activation Visualization [35 pts]

In this part, we will investigate a method to visualize the activation map of an image through a classification network and show intuitively how each part of the image contribute to the prediction. We have provided some starter code in `part2.py`.

To observe a meaningful pattern, we construct a custom dataset that localizes the Fashion-MNIST image with the help of MNIST images. Each image in this new dataset will be a 2×2 grid of one Fashion-MNIST image and three MNIST images. The Fashion-MNIST image is placed at a random grid, where the other three grids will be randomly-chosen MNIST images. This part is implemented for you in class `GridDataset`.

To apply the visualization method, our network needs to contain a global average pooling (GAP) layer followed by a linear layer at the end. When visualizing, we replace GAP and linear layers with a 1×1 convolution layer using weights from the linear layer. Instead of C class scores, the network output is now

C 2D arrays corresponding to each of the C classes. If we plot them as heatmaps as shown in the figure below, we should see that at ground truth class, activation is higher at the position of the Fashion-MNIST image in the input image, implying that our model has learned to "look at" only the Fashion-MNIST images for classification.

Notes on dimensions:

1. A global average pooling layer reduces each $H \times W$ channel to a single value by simply taking the average of all HW values.
2. Suppose the input to GAP layer in the original network has shape (F, H, W) , it will become $(F, 1, 1)$ after GAP layer, so the linear layer has weight of shape (F, C) . In the adapted network, the 1×1 conv layer has F input channels and C output channels, and therefore has weight of shape $(F, C, 1, 1)$. Since linear layer and 1×1 conv layer have weights of the same size, we can transfer weights from the former to the latter with a simple reshaping.
3. In PyTorch, the shape of a data tensor always has a dimension for batch size N , which is the first dimension.




Figure 3: Example image from our custom dataset and its per-class activation maps

Complete the following:


1. **Report** the detailed architecture of your `self.base` module. Include information on hyperparameters chosen for training, and the accuracy on the test set. To make the visualization look nice, you should achieve over 80% on the test set. (10 pts)
2. Choose a correctly classified image from the evaluation on test set. **Report** its index in the test set and include plots of both the image and the activation maps of all classes. (10 pts)
3. **Submit** a program which contains your best combination of `self.base` module, optimizer and training parameters, along with the code to select a correctly classified image and to visualize the results. (15 pts)

Hints: Apart from the last two layers, your network should look similar to the one in the previous part, and takes a similar amount of time to train. Again, any test accuracy over **50%** will receive partial credits.

3 Semantic Segmentation [30 pts]

Besides image classification, Convolutional Neural Networks can also generate dense predictions. A popular application is semantic segmentation. In this part, you will design and implement your Convolutional Neural Networks to perform semantic segmentation on the Mini Facade dataset.

Mini Facade dataset consists of images of different cities around the world and diverse architectural styles (in .jpg format), shown as the image on the left. It also contains semantic segmentation labels (in .png format) in 5 different classes: `balcony`, `window`, `pillar`, `facade` and `others`. Your task is to train a network to convert image on the left to the labels on the right.



Note: The label images are plotted using 8-bit indexed color with pixel value listed in Table 1. We have provided you a visualization function using a “jet” color map.

Table 1: Label image consists of 5 classes, represented in [0, 4].

| class | color | pixel value |
|---------|--------|-------------|
| others | black | 0 |
| facade | blue | 1 |
| pillar | green | 2 |
| window | orange | 3 |
| balcony | red | 4 |

We have provided some starter code in `part3/train.py` which contains a dummy network. It uses a 1x1 convolution to convert 3 channels (RGB) to 5 channels, i.e. 5 heatmaps for each class, with cross-entropy loss. Your network should give an output of the same shape, but contains more different kinds of layers. In addition to layers in a classification

We will evaluate your model on its average precision (AP) on the test set (higher the better). Weve already provided you the code to evaluate AP and you can directly use it. An introduction to AP is available at https://scikit-learn.org/stable/modules/model_evaluation.html#precision-recall-f-measure-metrics.

Complete the following:

1. **Report** the detailed architecture of your model. Include information on hyperparameters chosen for training and a plot showing both training and validation loss across iterations. (10 pts)
2. Report the average precision on the test set. You can use provided function to calculate AP on the test set. You should only evaluate your model on the test set once. All hyperparameter tuning should be done on the validation set. We expect you to to achieve **0.45** AP on the test set. (10 pts)

3. **Submit** a program which contains your best combination of `self.base` module, optimizer and training parameters. (5 pts)
4. Select or take a photo of any building, preprocess it as you like and input it to your best trained model. Plot the output labels and comment qualitatively on why it works or doesn't work. **Submit** the image as `input.jpg` and the output labels as `output.png`. (5 pts) Like HW1 Part3, we will have a contest where every student can vote for the best-labeled images. The winner will get **Extra Credits**.(2 pts)

Hints: You should still choose layers from classes under `torch.nn`. In addition to layers mentioned in the previous part, you might want to try `nn.Upsample` that upsamples the activation map by a simple interpolation, and `nn.ConvTranspose2d` that upsamples by performing transpose convolution. Any test AP over **0.25** will receive partial credits. Below are some architectures from papers that you can try:

- Jonathan Long, Evan Shelhamer, Trevor Darrell. Fully Convolutional Networks for Semantic Segmentation. CVPR 2015.
- Olaf Ronneberger, Philipp Fischer, Thomas Brox. U-Net: Convolutional Networks for Biomedical Image Segmentation. MICCAI 2015.
- Alejandro Newell, Kaiyu Yang, Jia Deng. Stacked Hourglass Networks for Human Pose Estimation. ECCV 2016.

You will easily get a high test AP if you use one of these models, but they can take a long time to train (hours with a CPU). Consider doing this part with the help of a GPU on Google Colab.

References

- [1] Jonathan Long, Evan Shelhamer, Trevor Darrell. Fully Convolutional Networks for Semantic Segmentation. CVPR 2015.
- [2] Han Xiao, Kashif Rasul, Roland Vollgraf. Fashion-MNIST: a Novel Image Dataset for Benchmarking Machine Learning Algorithms. arXiv:1708.07747.
- [3] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, Alexei A. Efros. Image-to-Image Translation with Conditional Adversarial Nets. CVPR 2017.
- [4] Radim Tyleček and Radim Šára. Spatial Pattern Templates for Recognition of Objects with Regular Structure. GCPR 2013.

Acknowledgement

The Mini Facade dataset are modified from CMP Facade Database by Radim Tyleček and Radim Šára. The problems are prepared by Dr. David Fouhey. We are grateful for their generosity.